

Analyze_ab_test_results_notebook

September 17, 2020

0.1 Analyze A/B Test Results

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the [RUBRIC](#).

Part I - Probability

To get started, let's import our libraries.

```
In [1]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
from statsmodels.stats.outliers_influence import variance_inflation_factor
from patsy import dmatrices

%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. **Use your dataframe to answer the questions in Quiz 1 of the classroom.**

a. Read in the dataset and take a look at the top few rows here:

```
In [2]: df = pd.read_csv('ab_data.csv')
df.head()
```

```
Out[2]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0

1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

b. Use the cell below to find the number of rows in the dataset.

```
In [3]: df.shape[0]
```

```
Out[3]: 294478
```

c. The number of unique users in the dataset.

```
In [4]: df.nunique().user_id
```

```
Out[4]: 290584
```

d. The proportion of users converted.

```
In [5]: df.converted.mean()
```

```
Out[5]: 0.11965919355605512
```

e. The number of times the new_page and treatment don't match.

```
In [6]: df_treatment_old = df.query('group == "treatment" and landing_page == "old_page"')
df_control_new = df.query('group == "control" and landing_page == "new_page"')
pd.concat([df_treatment_old, df_control_new]).user_id.count()
```

```
Out[6]: 3893
```

f. Do any of the rows have missing values?

```
In [7]: df.isna().sum()
```

```
Out[7]: user_id      0
timestamp    0
group        0
landing_page  0
converted    0
dtype: int64
```

Here I can see it doesn't have any missing value.

2. For the rows where **treatment** does not match with **new_page** or **control** does not match with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to figure out how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
In [8]: df2 = df.drop(df_treatment_old.index)
df2 = df2.drop(df_control_new.index)
```

```
In [9]: # Double Check all of the correct rows were removed - this should be 0
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].sha
```

```
Out[9]: 0
```

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_ids** are in **df2**?

```
In [10]: df2.user_id.nunique()
```

```
Out[10]: 290584
```

b. There is one **user_id** repeated in **df2**. What is it?

```
In [11]: df2[df2.user_id.duplicated() == True]
```

```
Out[11]:
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

c. What is the row information for the repeat **user_id**?

```
In [12]: df2[df2.user_id.duplicated() == True]
```

```
Out[12]:
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```
In [13]: df2.drop(df2[df2.user_id.duplicated() == True].index, inplace=True)
```

4. Use **df2** in the cells below to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
In [14]: df2.converted.mean()
```

```
Out[14]: 0.11959708724499628
```

b. Given that an individual was in the control group, what is the probability they converted?

```
In [15]: df2.query('group == "control"').converted.mean()
```

```
Out[15]: 0.1203863045004612
```

c. Given that an individual was in the treatment group, what is the probability they converted?

```
In [16]: df2.query('group == "treatment"').converted.mean()
```

```
Out[16]: 0.11880806551510564
```

d. What is the probability that an individual received the new page?

```
In [17]: df2.query('landing_page == "new_page").user_id.count() / df2.shape[0]
```

```
Out[17]: 0.50006194422266881
```

e. Consider your results from parts (a) through (d) above, and explain below whether you think there is sufficient evidence to conclude that the new treatment page leads to more conversions.

Answer : No, the treatment group has a less probability than the control group. Therefore, there is no evidence to conclude that the new treatment page leads to more conversions.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new} , which are the converted rates for the old and new pages.

Put your answer here.

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **conversion rate** for p_{new} under the null?

```
In [18]: df2.converted.mean()
```

```
Out[18]: 0.11959708724499628
```

b. What is the **conversion rate** for p_{old} under the null?

```
In [19]: df2['converted'].mean()
```

```
Out[19]: 0.11959708724499628
```

c. What is n_{new} , the number of individuals in the treatment group?

```
In [20]: df2.query('group == "treatment"').user_id.count()
```

```
Out[20]: 145310
```

d. What is n_{old} , the number of individuals in the control group?

```
In [21]: df2.query('group == "control"').user_id.count()
```

```
Out[21]: 145274
```

e. Simulate n_{new} transactions with a conversion rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
In [22]: df2['new_page_converted'] = df2.query('landing_page == "new_page"').converted
```

```
In [23]: df2.head()
```

```
Out[23]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	new_page_converted
0	NaN
1	NaN
2	0.0
3	0.0
4	NaN

f. Simulate n_{old} transactions with a conversion rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
In [24]: df2['old_page_converted'] = df2.query('landing_page == "old_page"').converted
```

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).


```
In [25]: observed_diff = df2.new_page_converted.mean() - df2.old_page_converted.mean()
```

h. Create 10,000 $p_{new} - p_{old}$ values using the same simulation process you used in parts (a) through (g) above. Store all 10,000 values in a NumPy array called **p_diffs**.

```
In [26]: diffs_list = []
sample_size = df2.shape[0]
for _ in range(10000):
    sample_df = df2.sample(sample_size, replace=True)
    pro_new = sample_df.query('landing_page == "new_page"').converted.mean()
    pro_old = sample_df.query('landing_page == "old_page"').converted.mean()
    diff = pro_new - pro_old
    diffs_list.append(diff)
p_diffs = np.array(diffs_list)
```

- i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
In [27]: plt.hist(p_diffs, facecolor='g', alpha=0.75);
```




- j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

```
In [28]: (p_diffs > observed_diff).mean()
```

```
Out[28]: 0.5040999999999999
```

```
In [29]: null_mean = 0
         null_vals = np.random.normal(null_mean, p_diffs.std(), sample_size)
         plt.hist(null_vals);

         plt.axvline(x=observed_diff, color = 'red');
```



```
In [30]: p_val = (null_vals > observed_diff).mean()
         p_val
```

```
Out[30]: 0.90486399801778483
```

- k. Please explain using the vocabulary you've learned in this course what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

Answer : It is the p-value. This p-value is greater than 0.05 so that we cannot reject the null hypothesis. We can conclude there is not difference between the new and old pages

- l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer to the number of rows associated with the old page and new pages, respectively.

```
In [31]: #import warnings
         #warnings.simplefilter("ignore")
         import statsmodels.api as sm

         convert_old = df2.query('landing_page == "old_page").converted.sum()
         convert_new = df2.query('landing_page == "new_page").converted.sum()
         n_old = df2.query('landing_page == "old_page").user_id.count()
         n_new = df2.query('landing_page == "new_page").user_id.count()
```

```
/opt/conda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas
from pandas.core import datetools
```

```
In [32]: df2.query('landing_page == "old_page").converted.sum(),
df2.query('landing_page == "old_page").user_id.count()
```

```
Out[32]: 145274
```

- m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. [Here](#) is a helpful link on using the built in.

```
In [33]: count = np.array([convert_new, convert_old])
nobs = np.array([n_new, n_old])

z_score, p_val_z = sm.stats.proportions_ztest(count, nobs, alternative='larger')
z_score, p_val_z
```

```
Out[33]: (-1.3109241984234394, 0.90505831275902449)
```

- n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

Answer : z-score is -1.310924 and the p-value is 0.905058 which is the same in parts j.

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

- a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Answer : Logistic regression

- b. The goal is to use **statsmodels** to fit the regression model you specified in part a. to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create in `df2` a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
In [34]: df2['intercept'] = 1
df2[['ab_page', 'ab_page_temp']] = pd.get_dummies(df2.landing_page)
df2.head()
```

```
Out[34]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	

4	864975	2017-01-21 01:52:26.210827	control	old_page	1
---	--------	----------------------------	---------	----------	---

	new_page_converted	old_page_converted	intercept	ab_page	ab_page_temp
0	NaN	0.0	1	0	1
1	NaN	0.0	1	0	1
2	0.0	NaN	1	1	0
3	0.0	NaN	1	1	0
4	NaN	1.0	1	0	1

```
In [35]: df2.drop('ab_page_temp', axis=1, inplace=True)
```

```
In [36]: df2.head()
```

```
Out [36]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	new_page_converted	old_page_converted	intercept	ab_page
0	NaN	0.0	1	0
1	NaN	0.0	1	0
2	0.0	NaN	1	1
3	0.0	NaN	1	1
4	NaN	1.0	1	0

- c. Use **statsmodels** to instantiate your regression model on the two columns you created in part b., then fit the model using the two columns you created in part b. to predict whether or not an individual converts.

```
In [38]: logistic_model = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
result = logistic_model.fit()
#result.summary()
result.summary2()
```

```
Optimization terminated successfully.
Current function value: 0.366118
Iterations 6
```

```
Out [38]: <class 'statsmodels.iolib.summary2.Summary'>
"""
Results: Logit
=====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable:    converted              Pseudo R-squared:    0.000
Date:                  2020-09-17 12:19      AIC:                  212780.3502
No. Observations:      290584                BIC:                  212801.5095
```

```

Df Model:          1          Log-Likelihood:   -1.0639e+05
Df Residuals:      290582      LL-Null:         -1.0639e+05
Converged:         1.0000      Scale:           1.0000
-----
                Coef.   Std.Err.    z      P>|z|    [0.025   0.975]
-----
intercept    -1.9888    0.0081  -246.6690  0.0000   -2.0046   -1.9730
ab_page      -0.0150    0.0114   -1.3109  0.1899   -0.0374    0.0074
=====
"""

```

- d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
In [39]: np.exp(result.params)
```

```
Out[39]: intercept    0.136863
         ab_page      0.985123
         dtype: float64
```

- e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint**: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in **Part II**?

Answer : - p-value associated with ab_page is 0.19. The hypothesis in Part II is the difference between converted rates for the new and old pages is zero. But the hypothesis in Part III is there is no effect on conversion rates with respect to the new and old pages.

- f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

Answer : - We should consider other factors into the regression model as they might influence the conversions too. - Factors like device on which tests were taken or course which was looked at, prior academic background, age, might alter experience and ultimately, conversions. These limitations should be kept in mind while making the final decision. - The disadvantages to adding additional terms into the regression model is that even with additional factors we can never account for all influencing factors or accomodate them. Plus, small pilots and pivots sometimes work better in practice than long-drawn research without execution.

- g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. [Here](#) are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables**. Provide the statistical output as well as a written response to answer this question.

```
In [40]: df_additional = pd.read_csv('countries.csv')
df_additional.head()
```

```
Out[40]:
```

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK

```
In [41]: df2 = df2.join(df_additional.set_index('user_id'), on='user_id')
```

```
In [42]: df2.head()
```

```
Out[42]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	new_page_converted	old_page_converted	intercept	ab_page	country
0	NaN	0.0	1	0	US
1	NaN	0.0	1	0	US
2	0.0	NaN	1	1	US
3	0.0	NaN	1	1	US
4	NaN	1.0	1	0	US

```
In [43]: df2[['CA', 'UK', 'US']] = pd.get_dummies(df2.country)
df2.head()
```

```
Out[43]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	new_page_converted	old_page_converted	intercept	ab_page	country	CA	UK	\
0	NaN	0.0	1	0	US	0	0	
1	NaN	0.0	1	0	US	0	0	
2	0.0	NaN	1	1	US	0	0	
3	0.0	NaN	1	1	US	0	0	
4	NaN	1.0	1	0	US	0	0	

	US
0	1
1	1
2	1

```
3    1
4    1
```

```
In [45]: logistic_model_add_country = sm.Logit(df2['converted'], df2[['intercept', 'ab_page', 'US', 'UK'])
result2 = logistic_model_add_country.fit()
#result.summary()
result.summary2()
```

```
Optimization terminated successfully.
Current function value: 0.366113
Iterations 6
```

```
Out[45]: <class 'statsmodels.iolib.summary2.Summary'>
"""
```

```

                        Results: Logit
=====
Model:                  Logit                  No. Iterations:    6.0000
Dependent Variable: converted                Pseudo R-squared: 0.000
Date:                   2020-09-17 12:20      AIC:                212780.3502
No. Observations:      290584                BIC:                212801.5095
Df Model:               1                    Log-Likelihood:   -1.0639e+05
Df Residuals:           290582                LL-Null:            -1.0639e+05
Converged:              1.0000                Scale:              1.0000
-----
                        Coef.   Std.Err.   z         P>|z|    [0.025   0.975]
-----
intercept   -1.9888    0.0081  -246.6690  0.0000   -2.0046  -1.9730
ab_page     -0.0150    0.0114   -1.3109  0.1899   -0.0374   0.0074
=====
"""
```

```
In [46]: np.exp(result2.params)
```

```
Out[46]: intercept    0.131332
ab_page    0.985168
US         1.041599
UK         1.051944
dtype: float64
```

```
In [47]: y, X = dmatrices('converted ~ ab_page + US + UK', df2, return_type = 'dataframe')

vif = pd.DataFrame()
vif["VIF Factor"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif["features"] = X.columns
vif
```

```
Out[47]:   VIF Factor  features
0    21.055923  Intercept
```

1	1.000006	ab_page
2	4.502223	US
3	4.502236	UK

The effects between countries are almost same.

- h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

Results for logistic regression models shows that none of the variable, treatment or country is significant in predicting conversion rate. The ecommerce company should really consider either running the experiment longer or use their domain knowledge to identify some possible key characteristic of their best customers to help predict conversion rate.

Finishing Up

Congratulations! You have reached the end of the A/B Test Results project! You should be very proud of all you have accomplished!

Tip: Once you are satisfied with your work here, check over your report to make sure that it is satisfies all the areas of the rubric (found on the project submission page at the end of the lesson). You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

0.2 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File > Download as** sub-menu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [48]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```

```
Out[48]: 0
```